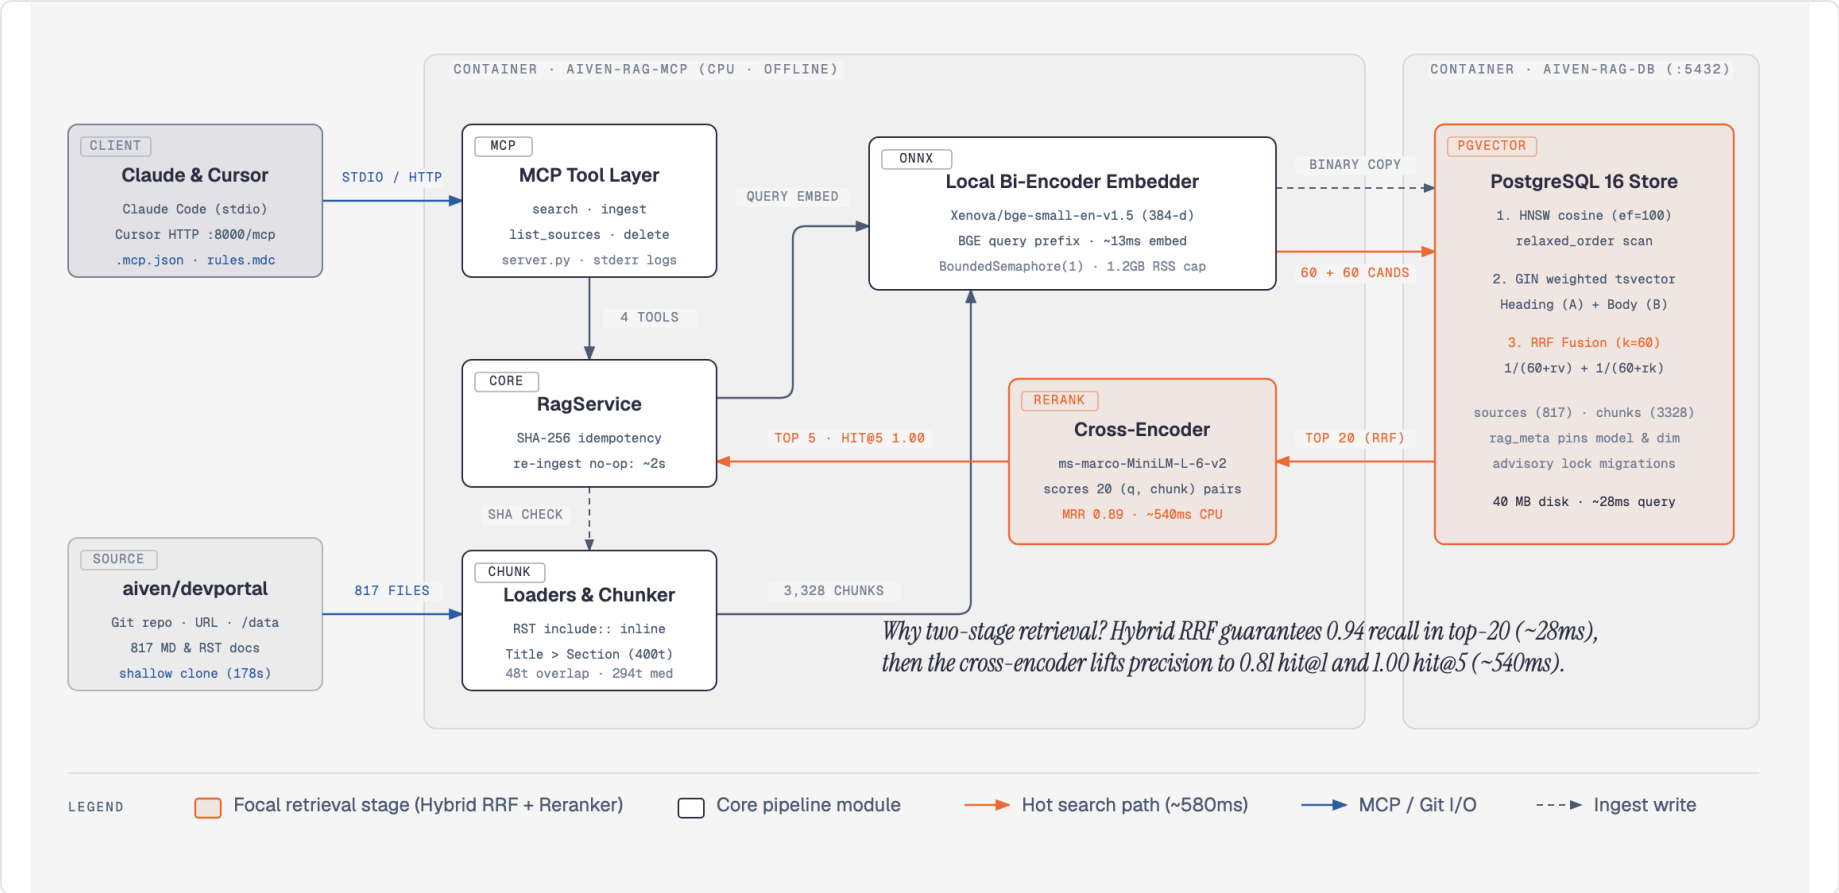
1. The solution: a local RAG server any MCP client can use



- **One command:** `docker compose up -d --build`, then `rag-mcp ingest https://github.com/aiven/devportal`.
- **No API keys and no runtime network:** local ONNX models (`bge-small-en-v1.5` + `ms-marco-MiniLM-L-6-v2`) are baked into the image, data lives on a named volume, and migrations run under a Postgres advisory lock with model/dimension pinning.
- **Idempotent ingest & Cursor-ready:** SHA-256 per file (atomic replace, `--prune`), plus `.cursor/mcp.json`, rules, and `AGENTS.md`.

2. Load: measured on a 4 vCPU / 16 GB VM

Ingest (full aiven/devportal)	
817 files → 3,328 chunks (median 294 tokens)	**178 s** (≈19 chunks/s, CPU embedding)
Re-ingest, nothing changed (clone + hash)	**≈2 s**
Database size	40 MB (chunks 21 MB + indexes 10 MB)

Search through a real MCP client (HTTP)	**conc 1**	**conc 4**	**conc 16**	**max req/s**	**server RSS**
hybrid + rerank (default) p50 / p95	603 / 720 ms	2.4 / 2.5 s	9.8 / 10.8 s	**1.6**	1.2 GB
hybrid, reranker off p50 / p95	30 / 51 ms	78 / 131 ms	321 / 552 ms	**46**	0.3 GB

Retrieval quality (32 labeled questions)	**hit@1**	**hit@5**	**MRR@10**
vector / keyword / hybrid	0.75 / 0.31 / 0.72	0.91 / 0.72 / 0.94	0.83 / 0.48 / 0.82
hybrid + rerank	**0.81**	**1.00**	**0.89**

pgvector at scale (synthetic, vector part only)	**100k**	**300k**	**1M chunks**
HNSW build / index size	6 s / 195 MB	18 s / 586 MB	**276 s** / 1.95 GB
kNN p50 / p95 (ef_search 100)	1.1 / 1.9 ms	1.9 / 5.2 ms	1.6 / 5.2 ms

Raw data is in `docs/measurements/`; reproduce with `scripts/benchmark.py` and `scripts/scale_test.py`. The scale test's recall@10 was 1.0, but it is measured on clustered synthetic data, so no recall claim is made from it.

3. Where it breaks

1. **The CPU reranker is the capacity ceiling.** One rerank of 20 pairs uses every core for about 550 ms, so a container does about 1.6 searches/s. Above that, requests queue: 16 concurrent clients wait about 10 s each.
2. **Unbounded inference ran the host out of memory (found by the load test, now fixed).** Each in-flight ONNX run keeps its own activation buffers. 16 parallel searches took the server to **9.8 GB RSS**, and a second process doing the same pushed the 16 GB VM into a memory stall lasting about 50 minutes. A per-model semaphore now caps it at **1.2 GB** with the same throughput.
3. **HNSW past RAM.** At 1M chunks the graph (1.95 GB) outgrows `maintenance_work_mem`, so the build goes from 18 s to 276 s. Queries stay at about 2 ms only while the index fits in memory; past RAM every query reads from disk. Docker's default 64 MB `/dev/shm` also breaks parallel builds, so compose sets 1 GB.
4. **Long synchronous ingest.** A full devportal ingest is one 3-minute tool call. Clients with short tool timeouts give up, though the ingest still finishes and progress is reported.
5. **One process per stdio session.** Each Claude window that runs `docker exec ... serve` loads its own models (≈0.8–1.2 GB), so ten windows need about 10 GB. Cursor shares the single HTTP process.
6. **Retrieval quality limits.** The model is English-only with a 512-token window. Long RST tables split across chunks lose their header row. Keyword-only search is weak on paraphrased questions (hit@5 0.72).

4. The next level

Problem	Next step
CPU-bound rerank, queueing	Move embedding and rerank to an inference service (e.g. TEI) with dynamic batching or a GPU. Shed load past a queue limit (return unreranked results instead of waiting 10 s). Cache query embeddings and results.
One process per session	One shared HTTP MCP server for all clients, with OAuth (MCP authorization spec). stdio becomes a thin proxy.
Long ingest	Async ingest jobs (`ingest` returns a job id, plus an `ingest_status` tool). A worker pool. Git-diff or webhook sync, so only changed files are fetched.
Index growth	`halfvec` (half the memory), or binary quantization with a re-score step. Partition by collection. Past about 10M chunks, use read replicas or a dedicated vector store.
Quality ceiling	Table-aware chunking (repeat the header row per chunk). A larger model (bge-base / e5) behind a re-index migration. Query rewriting for vague questions. A bigger eval set built from real queries.
Multi-tenant use	Per-tenant collections with row-level security, plus a quota per tool.

5. Finding and monitoring problems

Already in place - A structured log line per search with stage timings and the top score: `search mode=hybrid reranked=True query_chars=60 results=5 top_score=6.7001 embed_ms=13.3 retrieve_ms=28.2 rerank_ms=542.7 total_ms=584.2`. That one line shows the reranker is 93% of the latency. The query text is not logged. - The ingest report lists `added / updated / unchanged / pruned / failed` with a reason per failed file. One bad file never aborts the batch. - `rag-mcp eval` gives hit@k, MRR and answer-in-top-5 per mode. `benchmark.py` and `scale_test.py` are repeatable load and scale checks. - Compose healthchecks, plus readable tool errors such as “Cannot connect to PostgreSQL”.

Next to add (OpenTelemetry) - Metrics: p50/p95 per stage, wait time at the inference gate, rate of zero-result and low-top-score searches (the docs probably don’t cover the question), ingest failures, DB pool saturation, HNSW index size vs RAM. - Alerts: p95 > 2 s, tool error rate > 1%, any ingest failure, queue wait > 1 s. - Quality: run the eval in CI as a regression gate on every retrieval change. Each week, sample low-score queries (with consent), label them, and add them to the eval set.

6. Three tradeoffs

1. A cross-encoder reranker by default. hit@1 goes from 0.72 to 0.81, hit@5 from 0.94 to 1.00, and MRR from 0.82 to 0.89. The cost is **20× latency (30 → 600 ms) and 28× lower throughput**. I chose it because an agent's latency budget is seconds and wrong context costs more than half a second. It can be turned off with one env var, and the candidate count is tunable (10: MRR 0.86 at 280 ms; 30: 0.91 at 900 ms).

2. Hybrid (vector + full text, RRF) instead of vector only. Hybrid alone is *not* better at hit@1 (0.72 vs 0.75), but it finds more within the top 5 (0.94 vs 0.91) and is the safety net for exact identifiers (`max_connections`, CLI flags) that a small embedding model blurs. That gives the reranker better candidates. The cost is a second index, about 20 ms more and more moving parts. RRF fusion needs no score calibration.

3. A small local model baked into the image instead of a larger or hosted one. You get no key, no network, no per-query cost, a 3-minute full ingest on 4 CPU cores, and nothing leaves the machine. The cost is a lower quality ceiling: 384 dimensions, English only, and a 512-token window. The model name and dimension are pinned in the DB, so upgrading is an explicit re-index rather than silent corruption.